

IT Scripting and Automation

Unix/Linux Filesystem

Lecturer: Art Ó Coileáin

The Filesystem

- The file system is responsible for representing and organising the system's storage resources.
- The file system can be thought of as comprising four main components:
 - **A namespace:** a way to name things and organise them in a hierarchy.
 - **An API:** a set of system calls for navigating and manipulating objects.
 - **A security model:** a scheme for protecting, hiding, and sharing things.
 - **An implementation:** software to tie the logical model to the hardware.
- The file system is presented as a single unified hierarchy that starts at the directory **/** and continue downward through an arbitrary number of subdirectories.
- **/** is also called the root directory.

Pathnames

- The list of directories that must be traversed to locate a particular file plus that file's filename form a **pathname**.
- Pathnames can be:
 - **absolute** such as `'/tmp/foo'` or
 - **relative** such as `'bookN1/topic'`.
- Relative pathnames are interpreted **starting** at the current directory.
- Usually the terms: pathname and path are used interchangeably.

Limitations:

- The filesystem can be arbitrarily deep however, each component of a pathname (each directory name) must have no more than 255 characters long.
- There is also a limit on the path length, you can pass into the kernel as system call argument: 4095 bytes on Linux, 1023 bytes in older systems. To access a file with a pathname longer than this, you must `cd` to an intermediate directory and use relative pathnames.

The Naming of Files & Mount

- The naming of files and directories is basically unrestricted, except that:
 - names are limited in length and
 - must not contain slash characters or nulls.

Mount:

- Filesystem are attached to the 'file tree' with the command **mount**.
- Mount maps a directory with the existing 'file tree', called the mount point, to the root of the newly attached filesystem.
- **NB:** The previous contents of the mount point become inaccessible as long as another filesystem is mounted there.
- Mount points are usually empty directories.

Mount

Usage:

- `$ su -c mount /dev/sda4 /users`
- This command installs the filesystem stored on the disk partition represented by `/dev/sda4` under the path `/users`.
- `'ls'` command can be used to see the filesystem's content.
- To detach filesystems, the `umount` command is used :
`umount/mnt/cdrom`

File systems mounted at Boot-time:

- A list of filesystems that are customary mounted on a particular system is kept in the `/etc/fstab` or `/etc/filesystems` file. These filesystems are automatically mounted at boot time.

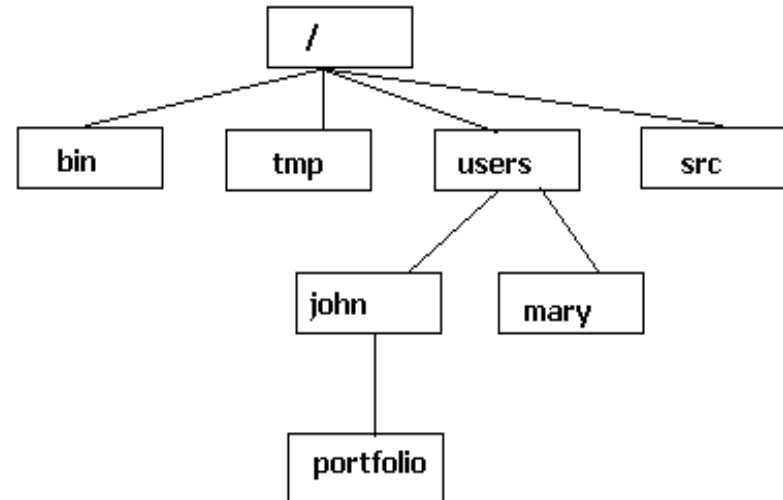
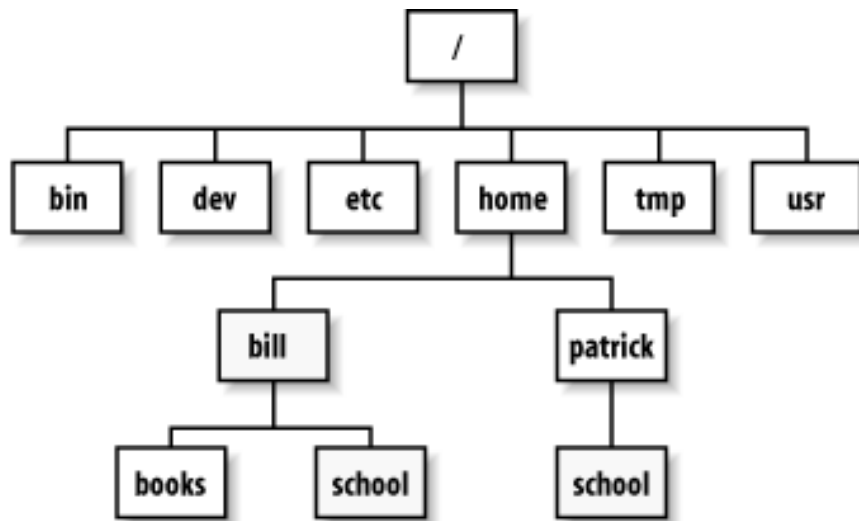
Organisation of the 'File Tree'

The main parts of the root filesystem are:

- **/etc** for critical systems and configuration files,
- **/sbin** holds commands needed to boot the system
- **/bin** for important utilities which are needed in single user mode
- **/tmp** for temporary files
- **/dev** contains special device files that control access to peripheral devices
- **/usr** where most of the standard programs are kept and most libraries. It is usually separate filesystem for convenience in administration
- **/lib** or **/usr/lib** keep shared library files
- **/var** for log files, accounting information, and various other items that grow or change rapidly and vary on each host. It is also a good idea to put it on its own filesystem
- **/home** or home directory of users are also kept on a separate filesystem
- **/proc** provides information about running processes and the kernel

Unix-like 'File trees'

- On many systems, the [hier](#) man page ([filesystem](#) on Solaris) outlines some general guidelines for the layout of the filesystem.



Disk partitions

AIX	<code>/dev/hdisk2</code> (refers to the entire disk)
FreeBSD	<code>/dev/da0s1e</code> (short form: <code>/dev/da1c</code>)
HP-UX	<code>dev/rdisk/c0t4d0</code>
Linux	<code>/dev/sdc1</code>
Solaris	<code>/dev/rdisk/c0t4d0s7</code>
Tru64	<code>/dev/rdisk/dsk2c</code>

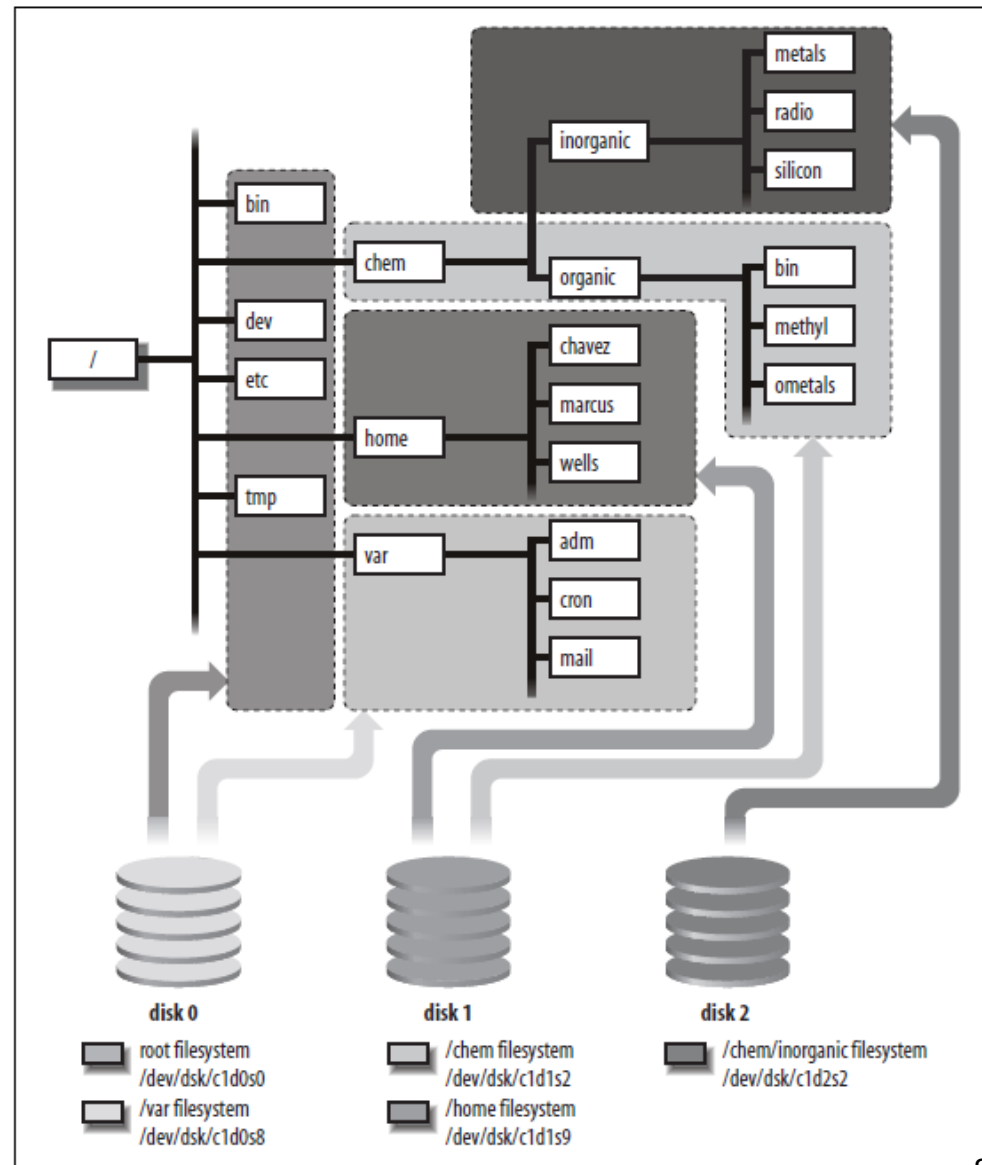
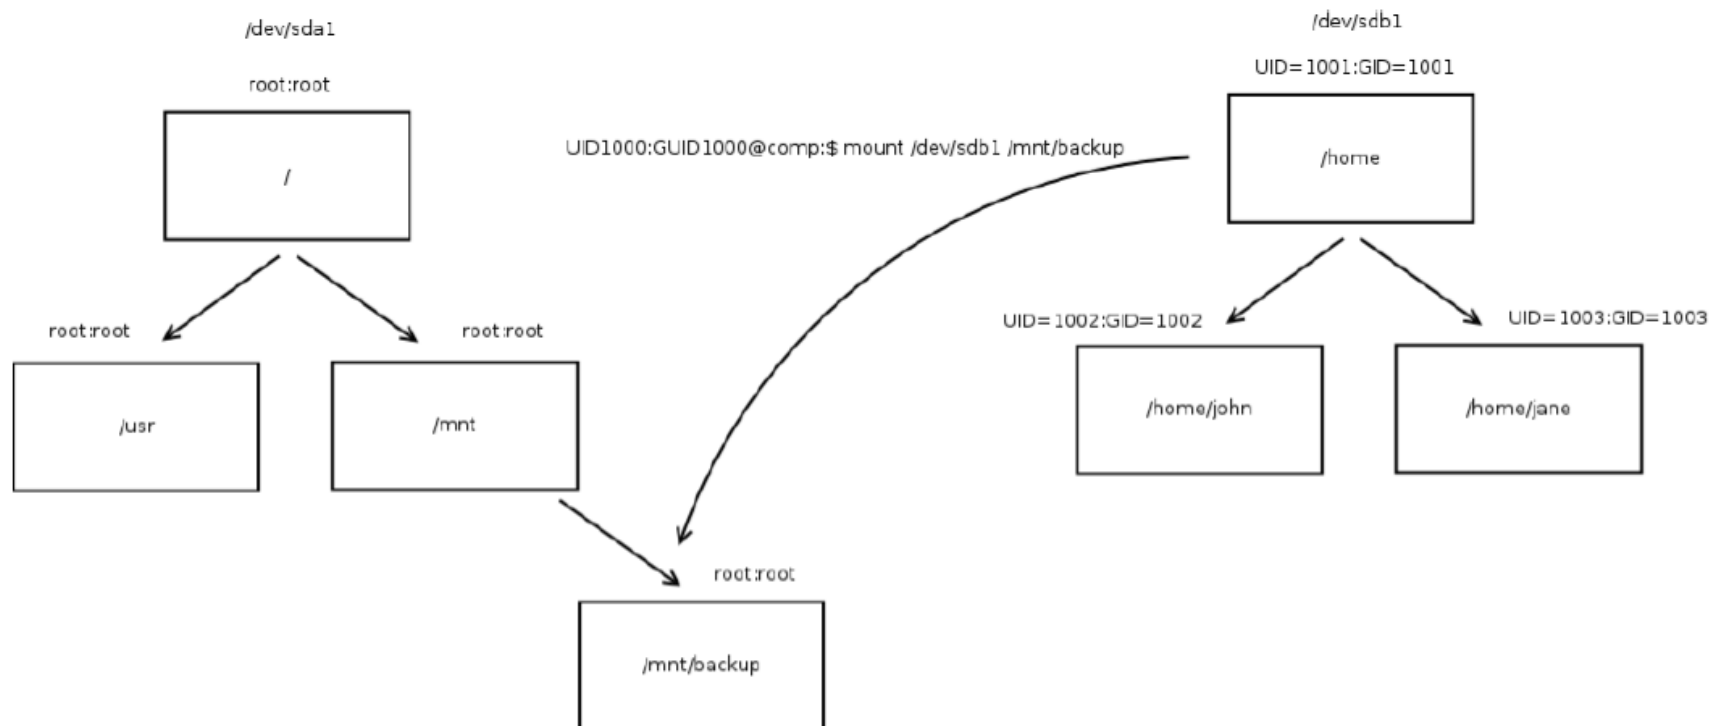


Figure 10-1. Mounting disk partitions within the Unix filesystem

Access to files on other partitions...

- Consider a system with 2 hard drives, having:
 - root filesystem at /dev/sda1 and
 - /home filesystem at /dev/sdb1.
- Now, a process with UID=1000:GID=1000 wants to mount /dev/sdb1 file system at /mnt/backup.



File Types

- The file type can be determined by using the command: `ls -ld` or `ls -l`

```
$ ls -ld /usr/include
```

```
dwxr-xr-x 27 root root 4096 Jul 15 20:57 /usr/include
```

File-type encoding used by ls

File type	Symbol	Created by	Removed by
Regular file	-	editors, cp , etc.	rm
Directory	d	mkdir	rmdir , rm -r
Character device file	c	mknod	rm
Block device file	b	mknod	rm
Local domain socket	s	socket(2)	rm
Named pipe	p	mknod	rm
Symbolic link	l	ln -s	rm

- It is a good idea to use `rm -i file.txt` when deleting files (to confirm the deletion)

Chmod: change permissions

- **chmod 711 myprogram.sh** gives all permissions to the owner and execute only permissions to everyone else.

Permission encoding for chmod

Octal	Binary	Perms	Octal	Binary	Perms
0	000	— — —	4	100	r — —
1	001	— — x	5	101	r — x
2	010	— w —	6	110	r w —
3	011	— w x	7	111	r w x

Chown and chgrp: change ownership/group

- The **chown** command changes a file's ownership, and the **chgrp** command changes its group ownership.
- To change a file's group, you must either be:
 - the **owner of the file** and belong to the group you are changing to or
 - be the **superuser**.

Recursive:

- Like **chmod**, **chown** and **chgrp** offer the recursive **-R** flag to change the setting of a directory and all the files underneath it.
 - `$ su -c 'chown -R seanfilename.txt'`
 - `$ su -c 'chgrp -R students2015 filename.txt'`
- **chown** command can change both the owner and group of a file at one with the syntax: **chown user:groupfile ...**
 - `$ su -c 'chown -R sean:students2015 filename.txt'`

Symbolic links

- A symbolic or soft link points to a file by name. When the kernel comes upon a symbolic link in the course of looking up a pathname, it redirects its attention to the pathname stored as the contents of the link.

Hard Vs Soft:

- The difference between hard links and symbolic links is that a hard link is a direct reference, whereas a symbolic link is a reference by name.

Usage:

- Symbolic links can be created with `ln -s` and be removed with `rm`.
- A symbolic link can contain either an absolute or a relative path.
 - `$su -c ln -s MyDirectory/secure.txt /tmp/secure.txt`
- It links `MyDirectory/secure.txt` to `/tmp/secure.txt` with. It creates the symbolic link `MyDirectory/secure.txt` with a target of `/tmp/secure.txt`